



Denodo MongoDB Custom Wrapper

Revision 20170303

NOTE

This document is confidential and proprietary of **Denodo Technologies**.
No part of this document may be reproduced in any form by any means without prior written authorization of **Denodo Technologies**.

Copyright © 2017
Denodo Technologies Proprietary and Confidential

CONTENTS

1 INTRODUCTION.....	4
2 WHAT IS MONGODB?.....	5
3 ARCHITECTURE AND FEATURES.....	6
4 USAGE.....	8
4.1 COMPATIBILITY.....	8
4.2 IMPORTING THE CUSTOM WRAPPER INTO VDP.....	8
4.3 CREATING A MONGODB DATA SOURCE.....	9
4.4 CREATING A BASE VIEW.....	10
4.5 CREATE BASE VIEW EXAMPLE.....	11
4.6 CREATE BASE VIEW USING FIELDS.....	14
4.7 CREATE BASE VIEW USING INTROSPECTION QUERY.....	15
4.8 CREATE BASE VIEW WITH SSL.....	16
4.9 QUERYING MONGODB COLLECTIONS.....	17
4.10 INSERT.....	17
4.11 UPDATE.....	18
4.12 DELETE.....	18
5 LIMITATIONS.....	19

1 INTRODUCTION

mongodb-customwrapper is a Virtual DataPort custom wrapper for querying MongoDB collections.

It bridges (to some extent) the gap between NoSQL and relational databases by establishing a predefined schema for output, thus enabling SQL queries on MongoDB.

2 WHAT IS MONGODB?

[MongoDB](#) is one of the most prominent so-called NoSQL databases, although it would be more accurate to say that is a non-relational database, or as defined on their website: "A document-oriented, scalable and high performance database", developed by 10gen.

In MongoDB we can forget about the table: we have **collections** which are groups of **documents** that do not have to share any common schema. So documents are to MongoDB something very similar to what tuples represent to relational databases, only without the schema restrictions associated with the latter.

Another important feature is the way in which MongoDB stores documents: **BSON** (Binary JSON). The binary part is hidden from us when we work with MongoDB, and thus we see and treat these data in JSON format, so our documents are in fact JSON objects, and its **fields** are the equivalent to the columns of the relational world. But as we do not need the documents in a collection to share a common schema, each JSON document can have just the fields it needs, and all documents in a collection do not have to represent entities of the same nature.

And you can also forget about both the SQL language and the JDBC API. MongoDB does not use SQL but its own query API, and consequently offers a specialized driver for Java that acts as a client and interacts with the database.

A **database** consists of one or more collections, the documents in those collections, and an optional set of security credentials for controlling access.

3 ARCHITECTURE AND FEATURES

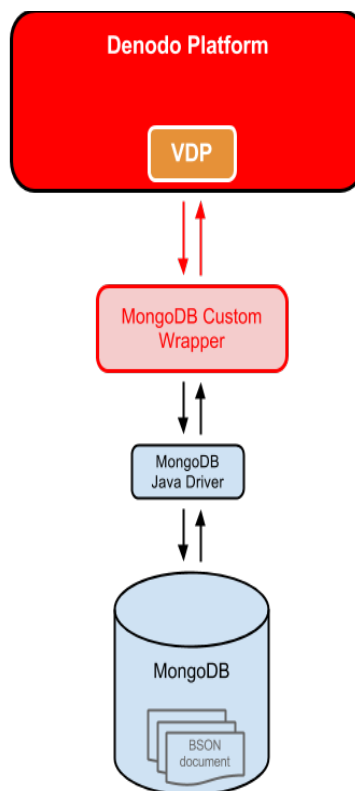
mongodb-customwrapper allows us to create base views on MongoDB collections and execute SQL queries on those collections.

This wrapper uses the official MongoDB Driver for Java.

This wrapper uses this driver to get a MongoClient object that represents a pool of connections to a MongoDB database. Optionally it can authenticate by user and password.

After making the connection, the custom wrapper can create base views to read collections and documents as if they were tables and rows. Each base view will be created on a specific collection.

Once the base views are created, VDP will be able to perform queries on MongoDB using the VQL Shell, create derived views, etc.



MongoDB Custom Wrapper Architecture

This is a brief summary of the wrapper's current features:

- User/password authentication is supported. User name and password can be specified as input values when building the base view.
- Base view result schema is defined as a comma-separated list of field names.
 - Extra fields in returned documents are ignored.
 - Lacking fields in returned documents are returned as null.
 - Types can be specified (from java.sql.Types).
 - Default is VARCHAR.
 - If documents in the same collection specify different types for fields with the same name, using VARCHAR (text in VDP) for that column will automatically perform the required conversions to show data from those documents.
- Alternatively the schema could be defined using an introspection query. **All documents** (remember, “schemaless”) retrieved by this query are analyzed to reveal their fields and build the base view schema. For this reason the query should retrieve a significant sample of the collection we are interested in.
 - Since common fields may hold different types of data the resulting structure is the highest common denominator between all the fields with the same name. In case of incompatible fields --such as integer and subdocuments-- VDP interprets them as of type text. But notice that **MongoDB is strict about types** and you must query for data using the correct type, so this fields would not be searchable.
- Create/update operations using simple fields are supported.
- Delete operations are supported.
- Field projections are delegated to MongoDB.
- Some query conditions are delegated to MongoDB.
 - AND and OR conditions.
 - Operators: =, <>, <, >, <=, >=, like, IS NULL, IS NOT NULL.
 - IS NULL operator gets the documents where the specified field has the null value or does not exist.
- ORDER BY clause is delegated to MongoDB.

4 USAGE

4.1 COMPATIBILITY

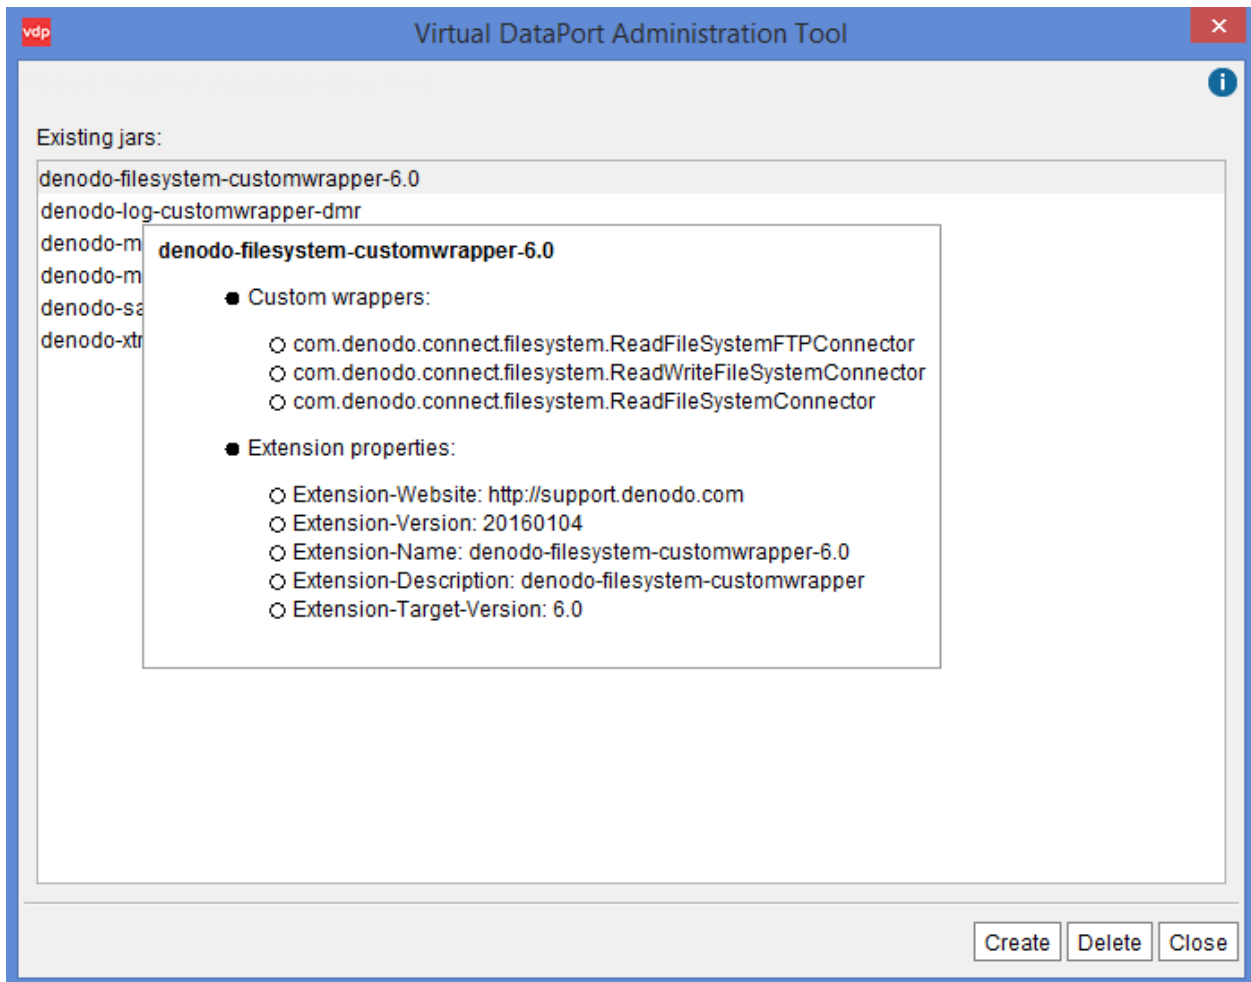
The driver does not support older versions of 2.4 MongoDB.

4.2 IMPORTING THE CUSTOM WRAPPER INTO VDP

In order to use the MongoDB Custom Wrapper in VDP, we must configure the Admin Tool to import the extension.

From the denodo-mongodb-customwrapper distribution, we will select the denodo-mongodb-customwrapper-\${version}-jar-with-dependencies.jar file and upload it to VDP.

No other jars are required as this one will already contain all the required dependencies, including the MongoDB Java driver classes.



MongoDB Extension in VDP

4.3 CREATING A MONGODB DATA SOURCE

Once the custom wrapper jar file has been uploaded to VDP using the Admin Tool, we can create new data sources for this custom wrapper --and their corresponding base views-- as usual.

Go to New → Data Source → Custom and specify the wrapper's class name `com.denodo.connect.mongodb.wrapper.MongoDBWrapper`. Also check 'Select Jars' and select the jar file of the custom wrapper.

The screenshot shows the Denodo configuration window for a MongoDB Data Source. The 'Connection' tab is selected. The 'Name' field is 'my_mongodb'. The 'Class name' is 'com.denodo.connect.mongodb.wrapper.MongoDBWrapper'. The 'Class path (optional)' field is empty, with a 'Browse' button next to it. The 'Select Jars' checkbox is checked. Below it, a list of JARs is displayed: 'denodo-log-customwrapper-dmr', 'denodo-mongodb-customwrapper-5.5-30', 'denodo-mongodb-customwrapper-6.0', 'denodo-salesforce-rest-wrapper-6.0', and 'denodo-xtrafuncs-vdp-date-dmr'.

MongoDB Data Source

4.4 CREATING A BASE VIEW

Once the custom wrapper has been registered, we will be asked by VDP to create a base view for it.

Base views created from the MongoDBWrapper need to fill in the Database parameter or the Connection string parameter, both of them are mutually exclusive. Moreover Collection is mandatory

- Database: Database name.
- Connection String: this parameter allows more options in the connection between the driver and mongodb. This is the pattern: [mongodb://]host1[:port1][,host2[:port2],...[,hostN[:portN]]]/database[?options]. The prefix mongodb:// is optional, the database is mandatory and user and password are not written in this field, but in the User and Password parameters
You can see all the options in the documentation of mongodb, in the following link about [Connection String URI Format](#).
If you introduce this parameter, host, port and database should be empty.
- Collection: Collection name that we will import as a table in VDP.

Four **optional** parameters:

- Host: Name of the computer or IP address where MongoDB is running, default is 127.0.0.1. Only with Database parameter.
- Port: Port number to connect to MongoDB, default is 27017. Only with Database

parameter.

- User/Password: Username and password to connect to MongoDB, if the authentication is enabled.

There are also two parameters that are **mutually exclusive**:

- Fields: The fields we would like to import as columns in VDP. We must keep the syntax `field1[:type1][,field2[:type2],...]`. Type, when specified, should be one of `java.sql.Types`. See a specific section below to learn more about the allowed syntax.
- Introspection query: Documents retrieved by this query will be analyzed to reveal their fields and build the view schema. An empty query selects all documents in the collection. This query requires **MongoDB syntax**.

4.5 CREATE BASE VIEW EXAMPLE

In the following example we want to import a product catalog collection of an E-Commerce site database.

At the beginning of the document, the schema must contain general product information, to facilitate searches of the entire catalog. Then, a details subdocument that contains fields that vary between product types.

As we are only interested in albums and films products we use the following introspection query:

```
{ type: { $in: [ 'Audio Album', 'Film' ] } }
```

Edit Wrapper Parameter values

Enter values for the following wrapper parameters:

Host	<input type="text" value="127.0.0.1"/>
Port	<input type="text" value="27017"/>
User	<input type="text"/>
Password	<input type="text"/>
Database	<input type="text" value="products"/>
Collection	<input type="text" value="product"/>
Connection String	<input type="text"/>
Fields	<input type="text"/>
Introspection query	<input type="text" value="{type:{<`\$in:[Audio Album','Film']}}"/>

Ok Cancel

MongoDB Base View edition using Host, Port and Database parameters

Edit Wrapper Parameter values

Enter values for the following wrapper parameters:

Host	<input type="text"/>
Port	<input type="text"/>
User	<input type="text"/>
Password	<input type="text"/>
Database	<input type="text"/>
Collection	<input type="text" value="product"/>
Connection String	<input type="text" value="mongodb://127.0.0.1:27017/products"/>
Fields	<input type="text"/>
Introspection query	<input type="text" value="{type:{<`\$in:[Audio Album','Film']}}"/>

Ok Cancel

MongoDB Base View edition using Connection String parameter

The resulting schema for this product catalog collection can be seen in the image below. It contains common product information like title, type, pricing and the details subdocument.

View schema		Metadata
View name: my_mongodb		
<input type="checkbox"/> PK	Field Name	Field Type
<input type="checkbox"/>	id_0	text
<input type="checkbox"/>	sku	text
<input type="checkbox"/>	type	text
<input type="checkbox"/>	title	text
<input type="checkbox"/>	description	text
<input type="checkbox"/>	publisher	text
<input type="checkbox"/>	pricing	my_mongodb_pricing
<input type="checkbox"/>	list	text
<input type="checkbox"/>	retail	text
<input type="checkbox"/>	details	my_mongodb_details

Set selected as PK

MongoDB Base View

Execute		Query Results
Results	Execution Trace	Stop Refresh Save
Total rows received: 2 (shown 2)		
id_0	sku	type
568bb83dd292501...	1000001	Audio Album
568bb83ed292501...	1000002	Audio Album

MongoDB Base View execution

Results		Execution Trace
Results	Execution Trace	Stop Refresh Save
Total rows received: 2 (shown 2)		
title	artist	genre
A Love Supreme [Original Recording Rei...	John Coltrane	Jazz

RESU... -> details

tracks

tracks_item

A Love Supreme Part I: Acknowledgement

A Love Supreme Part II - Resolution

A Love Supreme, Part III: Pursuance

A Love Supreme, Part IV-Psalm

Array and record details of MongoDB Base View execution

4.6 CREATE BASE VIEW USING FIELDS

The use of the Fields parameter for specifying the view's schema explicitly is recommended for base view creation when there are complex fields involved coming from mongoDB and there is a possibility that these data are not always returned in the same order or with the same type using an introspection query.

See an example of explicit schema specification (note braces for complex objects should be escaped):

```
_id : VARCHAR,  
firstName : VARCHAR,  
lastName : VARCHAR,  
age : VARCHAR,  
address: \{  
  city : VARCHAR,  
  state : VARCHAR,  
  streetAddress : VARCHAR,  
  state : VARCHAR,  
  postalCode : VARCHAR  
  \},  
phoneNumbers: array(\{  
  aaa : VARCHAR, type : VARCHAR,  
  number : VARCHAR  
  \}),  
cc:array(VARCHAR)
```

See how complex structures are specified using braces ({...}) and arrays using the array keyword and specifying the contents of the array between parentheses.

Also field names can be surrounded by single or double quotes (e.g. "phone number") if they contain white spaces or non-alphanumeric chars.

Edit Wrapper Parameter values

Enter values for the following wrapper parameters:

Host

127.0.0.1

Port

27017

User

Password

Database

test

Collection

address

Connection String

Fields

_id:VARCHAR,

firstName:VARCHAR

< >

Introspection query

Ok

Cancel

Mongo DB base View using Fields

Edit value of 'Fields'

_id:VARCHAR,

firstName:VARCHAR

,lastName:VARCHAR,

age:VARCHAR,

address:{city:VARCHAR,state:VARCHAR,streetAddress:VARCHAR,state:VARCHAR,postalCode:VARCHAR},

phoneNumbers:array({aaa:VARCHAR,type:VARCHAR,number:VARCHAR}),

cc:array(VARCHAR)

Fields Parameter

4.7 CREATE BASE VIEW USING INTROSPECTION QUERY

If you want to not define the schema, you should use the Introspection query parameter. For this you should use a query over the collection, following the syntax for the mongodb method `db.collection.find()`. You should escape the braces.

A example of introspection query is :

```
\{  "properties":  \{"LINEARID":  "1101121687137",  "observed":  "",
"COUNTYFP":  "051",  "RTTYP":  "M",  "FULLNAME":  "Airport Rd",  "MTFCC":
"S1400",  "STATEFP":  "35"\}\}
```

Edit Wrapper Parameter values

Enter values for the following wrapper parameters:

Host	<input type="text"/>
Port	<input type="text"/>
User	<input type="text"/>
Password	<input type="text"/>
Database	<input type="text"/>
Collection	<input type="text" value="roads"/>
Connection String	<input type="text" value="mongodb://127.0.0.1:27017/test"/>
Fields	<input type="text"/>
Introspection query	<input type="text" value='{"properties": {"LINEARD": 1101121687137}}'/>

Ok
Cancel

MongoDB Base View edition using Introspection query parameter

4.8 CREATE BASE VIEW WITH SSL

The MongoDB API for Java Driver allows the connection to a MongoDB Server that is set up with SSL. To create a Base view using SSL, have to add the parameter `ssl` with value `True` in the Connection String. SSL is one of the *Connection Options* of the *Connection String URI*.

In addition you have to import the certificate(used by MongoDB server) into the *TrustStore* of the Java Runtime Environment(JRE) of the platform. The section "Importing the Certificates of Data Sources (SSL Connections)" of the Installation Guide explains how to do this.

A example of the Base view with SSL:

Edit Wrapper Parameter values

Enter values for the following wrapper parameters:

Host	<input style="width: 80%;" type="text"/>
Port	<input style="width: 80%;" type="text"/>
User	<input style="width: 80%;" type="text"/>
Password	<input style="width: 80%;" type="text"/>
Database	<input style="width: 80%;" type="text"/>
Collection	<input style="width: 80%;" type="text" value="address"/>
Connection String	<input style="width: 80%;" type="text" value="localhost:27017/test?ssl=true"/>
Fields	<input style="width: 80%;" type="text"/>
Introspection query	<input style="width: 80%;" type="text"/>

MongoDB Base View using SSL

4.9 QUERYING MONGODB COLLECTIONS

From now on we can query using the VQL Shell tool of the Admin Tool, create derived views, etc.

Execute Quick Query Specify Where Expression Execute Query plan

Current sentence:

```
SELECT * FROM my_mongodb WHERE (details).genre = 'Jazz' CONTEXT ('i18n'='us_pst', 'cache_wait_for_load'='true')
```

Query Results Execution Trace

my_mongodb AS my_mongodb WHERE (details).genre = 'Jazz' CONTEXT ('i18n'='us_pst', 'cache_wait_for_load'='true') TRACE

Total rows received: 1 (shown 1)

id_0	sku	type	title	description	publisher	pricing	details
568bb83dd292501...	1000001	Audio Album	A Love Supreme	by John Coltrane	Sony Music	[Register]...	[Register]...




RESU... -> details			
title	artist	genre	tracks
A Love Supreme [Original Recording Rei...	John Coltrane	Jazz	[Array]...

Querying a MongoDB View

4.10 INSERT

This wrapper allows to make insertion with not complex fields



Execute Stop Load Clear Cut Copy Paste Previous Next

Database: admin Limit rows 150 Stop query when the limit is reached


```
INSERT INTO my_mongodb (id_0,city,pop,state)
VALUES (999999,'CORUNA','333','GL')
```

Output Execution Log

1 rows affected.

4.11 UPDATE

This wrapper allows to make updates in a document with not complex fields



Execute Stop Load Clear Cut Copy Paste Previous Next

Database: admin Limit rows 150 Stop query when the limit is reached


```
UPDATE my_mongodb
SET (city)=('Ferrol')
WHERE (id_0=999999);
```

Output Execution Log

1 rows affected.

4.12 DELETE

It is allowed to delete documents of a collection



Execute Stop Load Clear Cut Copy Paste Previous Next

Database: admin Limit rows 150 Stop query when the limit is reached

```
DELETE FROM my_mongodb
WHERE (id_0=999999);
```

Output Execution Log

1 rows affected.

5 LIMITATIONS

Base views

If you use the “Introspection Query” parameter when creating base views, the schema of the base view is inferred by VDP by doing a “best guess” based on the data returned at the time. So subsequent changes on the *real* schema or data types of the MongoDB-returned data might not be adequately detected by VDP.

This might affect views with complex types or arrays, and can be an issue if there is the possibility that changes in the underlying mongoDB data cause subfields in complex data be returned by mongoDB in different order or with different types/structure. In such cases, data could be returned in the wrong positions until a “Source Refresh” operation is performed.

It is therefore recommended to always explicitly set the schema using the “Fields” parameter when complex types or arrays are involved.

Write operations

Create/update operations with complex fields, arrays and registers, on the MongoDB server are not supported through this custom wrapper at its current version.