



# Denodo MongoDB Custom Wrapper

Revision 20180928

## NOTE

This document is confidential and proprietary of **Denodo Technologies**.  
No part of this document may be reproduced in any form by any means without prior written authorization of **Denodo Technologies**.

Copyright © 2018  
Denodo Technologies Proprietary and Confidential

## CONTENTS

<b>1 INTRODUCTION.....</b>	<b>4</b>
<b>2 WHAT IS MONGODB?.....</b>	<b>5</b>
<b>3 REQUIREMENTS.....</b>	<b>6</b>
<b>4 ARCHITECTURE AND FEATURES.....</b>	<b>7</b>
<b>5 USAGE.....</b>	<b>9</b>
5.1 IMPORTING THE CUSTOM WRAPPER INTO VDP.....	9
5.2 CREATING A MONGODB DATA SOURCE.....	10
5.3 CREATING A BASE VIEW.....	10
5.4 CREATE BASE VIEW EXAMPLE.....	11
5.5 CREATE BASE VIEW USING FIELDS.....	14
5.6 CREATE BASE VIEW USING INTROSPECTION QUERY.....	16
5.7 CREATE BASE VIEW WITH TLS/SSL.....	17
5.8 QUERYING MONGODB COLLECTIONS.....	18
5.9 INSERT.....	19
5.10 UPDATE.....	19
5.11 DELETE.....	19
5.12 USING MONGODB WRAPPER TO CONNECT TO AZURE COSMOS DB	20
<b>6 LIMITATIONS.....</b>	<b>24</b>
6.1 WRITE OPERATIONS.....	24
6.2 MONGODB BSON TIMESTAMP DATA TYPE.....	24
6.3 LIMITATIONS IN SELECT SENTENCES.....	24

## 1 INTRODUCTION

---

mongodb-customwrapper is a Virtual DataPort custom wrapper for querying MongoDB collections.

It bridges (to some extent) the gap between NoSQL and relational databases by establishing a predefined schema for output, thus enabling SQL queries on MongoDB.

## 2 WHAT IS MONGODB?

---

[MongoDB](#) is one of the most prominent so-called NoSQL databases, although it would be more accurate to say that is a non-relational database, or as defined on their website: "A document-oriented, scalable and high performance database", developed by 10gen.

In MongoDB we can forget about the table: we have **collections** which are groups of **documents** that do not have to share any common schema. So documents are to MongoDB something very similar to what tuples represent to relational databases, only without the schema restrictions associated with the latter.

Another important feature is the way in which MongoDB stores documents: **BSON** (Binary JSON). The binary part is hidden from us when we work with MongoDB, and thus we see and treat these data in JSON format, so our documents are in fact JSON objects, and its **fields** are the equivalent to the columns of the relational world. But as we do not need the documents in a collection to share a common schema, each JSON document can have just the fields it needs, and all documents in a collection do not have to represent entities of the same nature.

And you can also forget about both the SQL language and the JDBC API. MongoDB does not use SQL but its own query API, and consequently offers a specialized driver for Java that acts as a client and interacts with the database.

A **database** consists of one or more collections, the documents in those collections, and an optional set of security credentials for controlling access.

### 3 REQUIREMENTS

---

This Custom wrapper requires:

- When used with Denodo 5.5: At least the **Denodo update 20170529**
- When used with Denodo 6.0: At least the **Denodo update 20170515**
- When used with Denodo 7.0: No particular update is needed.
- MongoDB version 2.6 or newer.

## 4 ARCHITECTURE AND FEATURES

---

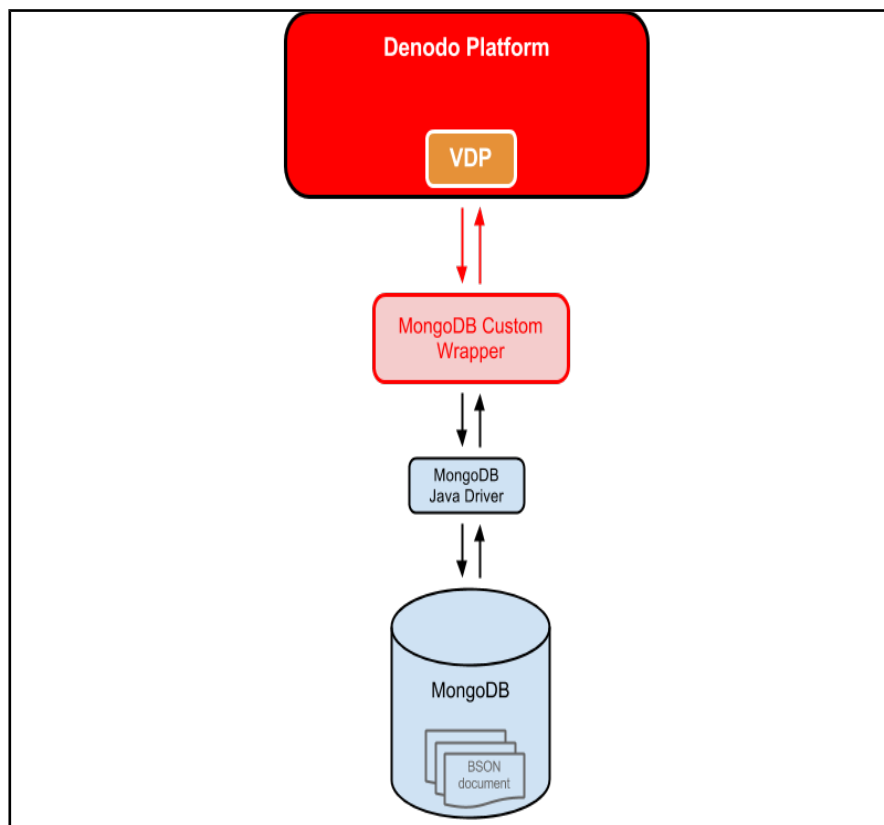
mongodb-customwrapper allows us to create base views on MongoDB collections and execute SQL queries on those collections.

This wrapper uses the official MongoDB Driver for Java.

It gets a MongoClient object that represents a pool of connections to a MongoDB database. Optionally it can authenticate by user and password.

After making the connection, the custom wrapper can create base views to read collections and documents as if they were tables and rows. Each base view will be created on a specific collection.

Once the base views are created, VDP will be able to perform queries on MongoDB using the VQL Shell, create derived views, etc.



MongoDB Custom Wrapper Architecture

This is a brief summary of the wrapper's current features:

- User/password authentication is supported. User and password can be specified as input values when building the base view.

- Base view result schema is defined as a comma-separated list of field names.
  - Extra fields in returned documents are ignored.
  - Lacking fields in returned documents are returned as null.
  - Types can be specified from constants in `java.sql.Types`, which correspond to SQL standard types.
    - Default is VARCHAR.
    - If documents in the same collection specify different types for fields with the same name, using VARCHAR (text in VDP) for that column will automatically perform the required conversions to show data from those documents.
    - Note that the SQL equivalent to MongoDB's "Date" type is SQL's TIMESTAMP, as MongoDB Date values contain both date and time.
    - MongoDB's internal "BSON Timestamp" data type is supported as SQL TIMESTAMP too, but with some limitations (see section at the end).
- Alternatively the schema could be defined using an introspection query. **All documents** (remember, "schemaless") retrieved by this query are analyzed to reveal their fields and build the base view schema. For this reason the query should retrieve a significant sample of the collection we are interested in.
  - Since common fields may hold different types of data the resulting structure is the highest common denominator between all the fields with the same name. In case of incompatible fields --such as integer and subdocuments-- VDP interprets them as of type text. But notice that **MongoDB is strict about types** and you must query for data using the correct type, so this fields would not be searchable.
- Create/update operations using simple fields are supported.
- Delete operations are supported.
- Field projections are delegated to MongoDB.
- Some query conditions are delegated to MongoDB.
  - AND and OR conditions.
  - Operators: =, <>, <, >, <=, >=, LIKE, IN, IS NULL, IS NOT NULL.
    - Note the IS NULL operator gets the documents where the specified field has the null value or does not exist.
- ORDER BY clause is delegated to MongoDB.

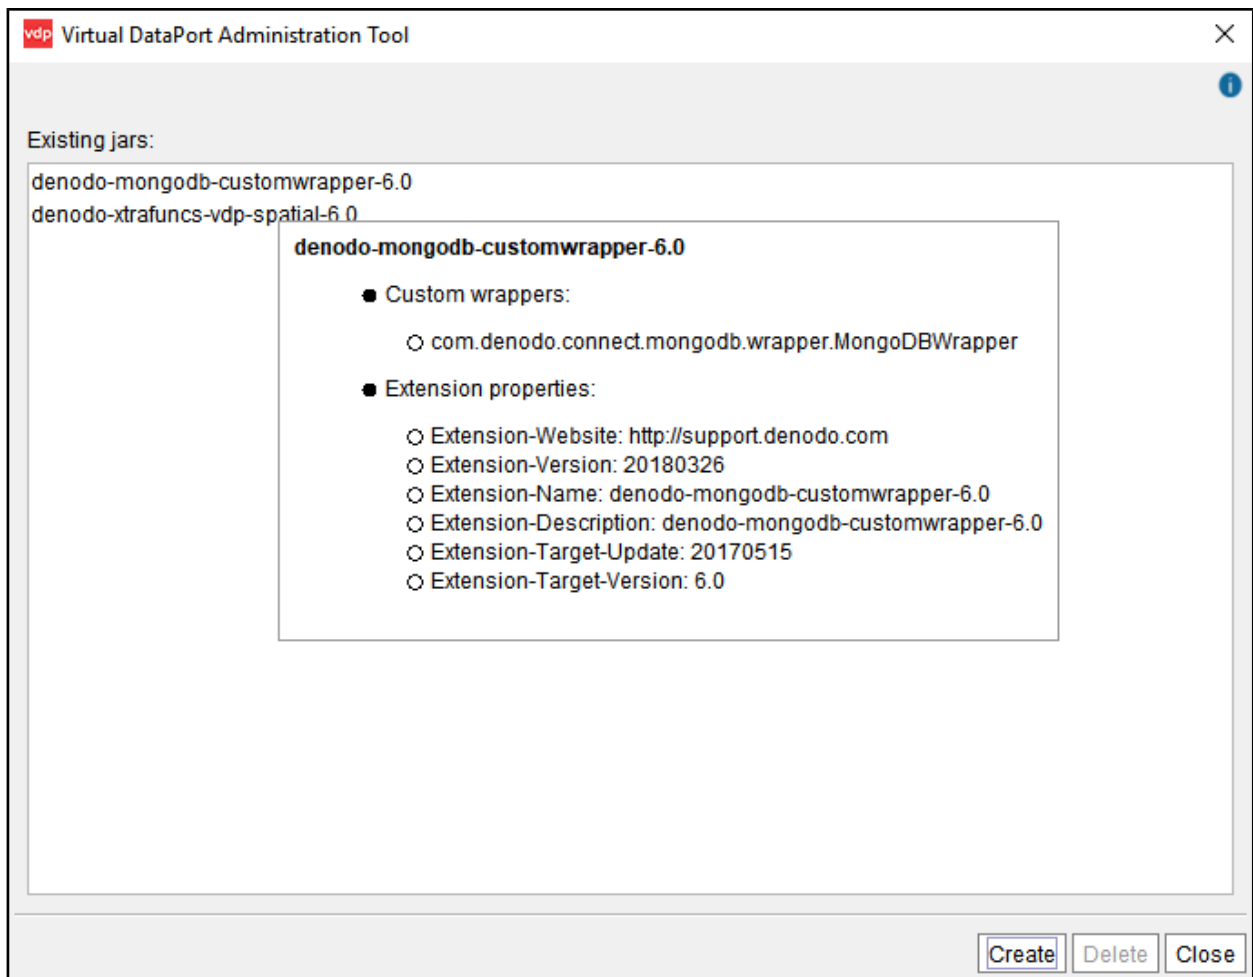
## 5 USAGE

### 5.1 IMPORTING THE CUSTOM WRAPPER INTO VDP

In order to use the MongoDB Custom Wrapper in VDP, we must configure the Admin Tool to import the extension.

From the denodo-mongodb-customwrapper distribution, we will select the denodo-mongodb-customwrapper-\${version}-jar-with-dependencies.jar file and upload it to VDP.

No other jars are required as this one will already contain all the required dependencies, including the MongoDB Java driver classes.



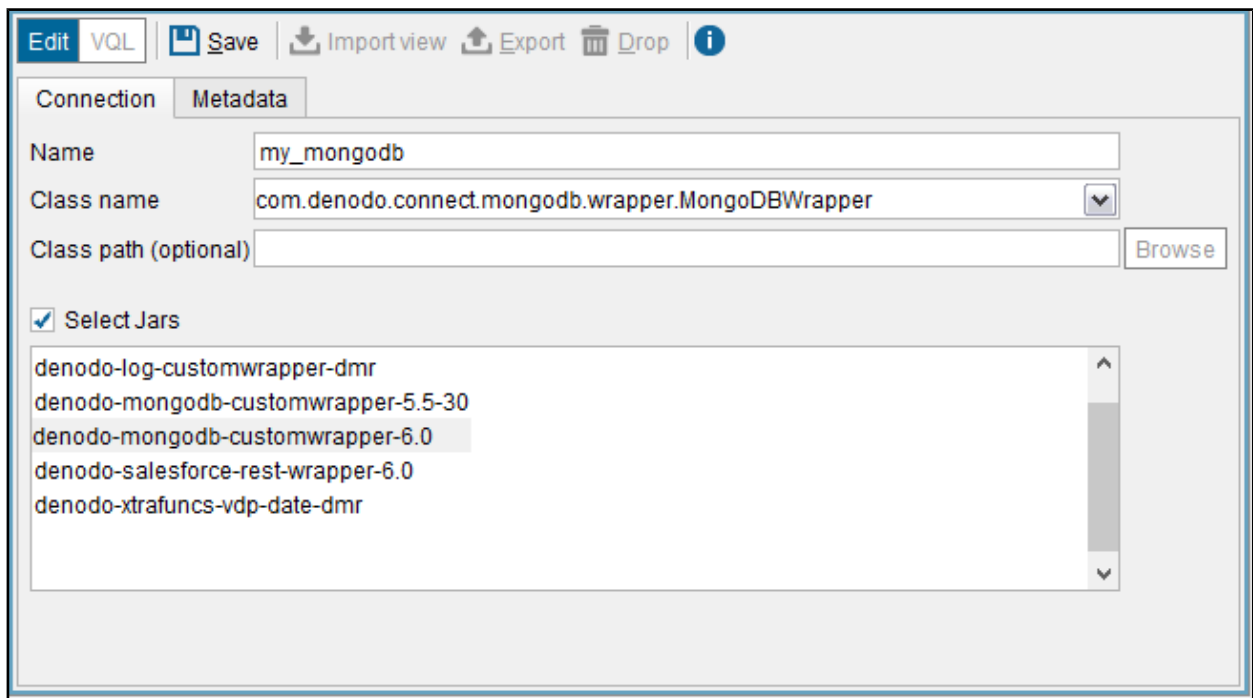
MongoDB Extension in VDP



## 5.2 CREATING A MONGODB DATA SOURCE

Once the custom wrapper jar file has been uploaded to VDP using the Admin Tool, we can create new data sources for this custom wrapper --and their corresponding base views-- as usual.

Go to New → Data Source → Custom and specify the wrapper's class name `com.denodo.connect.mongodb.wrapper.MongoDBWrapper`. Also check 'Select Jars' and select the jar file of the custom wrapper.



MongoDB Data Source

## 5.3 CREATING A BASE VIEW

Once the custom wrapper has been registered, we will be asked by VDP to create a base view for it.

Base views created from the `MongoDBWrapper` need to fill in the Database parameter or the Connection string parameter, both of them are mutually exclusive. Moreover Collection is mandatory

- **Database:** Database name.
- **Connection String:** this parameter allows more options in the connection between the driver and mongodb. This is the pattern:

```
[mongodb://]host1[:port1][,host2[:port2],...[,hostN[:portN]]]/database[?options]
```

The prefix `mongodb://` is optional, the database is mandatory and user and password are not written in this field, but in the User and Password parameters.

You can see all the options in the documentation of mongodb, in the following link about [Connection String URI Format](#).

If you introduce this parameter, Host, Port and Database parameters should be empty.

- **Collection:** Collection name that we will import as a table in VDP.

Five **optional** parameters:

- **Host:** Name of the computer or IP address where MongoDB is running, default is 127.0.0.1. Only with Database parameter.
- **Port:** Port number to connect to MongoDB, default is 27017. Only with Database parameter.
- **User/Password:** Username and password to connect to MongoDB, if the authentication is enabled.
- **Pass-through session credentials:** Enabling this checkbox allows VDP to execute the queries using the credentials of the user logged into VDP, so User and Password fields would be only used for introspection.
- **Use SSL/TLS:** Enabling this checkbox will make the wrapper use SSL/TLS. When using a Connection String, this can also be specified by means of the `"ssl=true"` option.

There are also two parameters that are **mutually exclusive**:

- **Fields:** The fields we would like to import as columns in VDP. We must keep the syntax

```
field1[:type1][,field2[:type2],...].
```

Type, when specified, should be one of the constants in `java.sql.Types` (note these are SQL standard types). See a specific section below to learn more about the allowed syntax.

- **Introspection query:** Documents retrieved by this query will be analyzed to reveal their fields and build the view schema. An empty query selects all documents in the collection, so, if the collection has a lot of documents, the creation of the base view can take a long time. In addition, it is recommendable to check the types of the new view inferred from the introspection. This query requires **MongoDB syntax**.

## 5.4 CREATE BASE VIEW EXAMPLE

In the following example we want to import a product catalog collection of an E-Commerce site database.

At the beginning of the document, the schema must contain general product information, to facilitate searches of the entire catalog. Then, a details subdocument that contains fields that vary between product types.

As we are only interested in albums and films products we use the following introspection query:

```
{ type: { $in: [ 'Audio Album', 'Film' ] } }
```

Edit Wrapper Parameter values

Enter values for the following wrapper parameters:

Host

127.0.0.1

Port

27017

☐

Pass-through session credentials

User

Password

Database

products

Collection

product

Connection String

☐

Use SSL/TLS

Fields

Introspection query

\{type:\{\$in:['Audio Album','Film']\}\}



Ok

Cancel

MongoDB Base View edition using Host, Port and Database parameters

### Edit Wrapper Parameter values

Enter values for the following wrapper parameters:

Host	<input type="text"/>
Port	<input type="text"/>
	<input type="checkbox"/> Pass-through session credentials
User	<input type="text"/>
Password	<input type="text"/>
Database	<input type="text"/>
Collection	<input type="text" value="product"/>
Connection String	<input type="text" value="mongodb://127.0.0.1:27017/products"/>
	<input type="checkbox"/> Use SSL/TLS
Fields	<input type="text"/> 
Introspection query	<input type="text" value="\{type:\{\$in:['Audio Album','Film']\}\}"/> 

Ok
Cancel

MongoDB Base View edition using Connection String parameter

The resulting schema for this product catalog collection can be seen in the image below. It contains common product information like title, type, pricing and the details subdocument.

View schema Metadata

View name: my\_mongodb

PK	Field Name	Field Type
<input type="checkbox"/>	id_0	text
<input type="checkbox"/>	sku	text
<input type="checkbox"/>	type	text
<input type="checkbox"/>	title	text
<input type="checkbox"/>	description	text
<input type="checkbox"/>	publisher	text
<input type="checkbox"/>	pricing	my_mongodb_pricing
<input type="checkbox"/>	list	text
<input type="checkbox"/>	retail	text
<input type="checkbox"/>	details	my_mongodb_details

Set selected as PK

MongoDB Base View

Execute Query Results

Results Execution Trace Stop Refresh Save db AS my\_mongodb WHERE type <> 'Book' CONTEXT ('i18n'='us\_pst', 'cache\_wait\_for\_load'='true') TRACE

Total rows received: 2 (shown 2)

id_0	sku	type	title	description	publisher	pricing	details
568bb83dd292501...	1000001	Audio Album	A Love Supreme	by John Coltrane	Sony Music	[Register]...	[Register]...
568bb83ed292501...	1000002	Audio Album	Love Song	by Khali Fong	Sony Music	[Register]...	[Register]...

MongoDB Base View execution

Results Execution Trace Stop Refresh Save Is FROM my\_mongodb AS my\_mongodb WHERE type <> 'Book' CONTEXT ('i18n'='us\_pst', 'cache\_wait\_for\_load'='true') TRACE

Total rows received: 2 (shown 2)

RESU... -> details

title	artist	genre	tracks
A Love Supreme [Original Recording Rei...	John Coltrane	Jazz	[Array]...

RESU... -> details -> tracks

tracks_item
A Love Supreme Part I: Acknowledgement
A Love Supreme Part II - Resolution
A Love Supreme, Part III: Pursuance
A Love Supreme, Part IV-Psalm

Array and record details of MongoDB Base View execution

## 5.5 CREATE BASE VIEW USING FIELDS

The use of the Fields parameter for specifying the view's schema explicitly is recommend for base view creation when there are complex fields involved coming from mongoDB and there is a possibility that these data are not always returned in the same order or with the same type using an introspection query.

See an example of explicit schema specification (note braces for complex objects should be escaped):

```
_id : VARCHAR,  
firstName : VARCHAR,  
lastName : VARCHAR,  
age : VARCHAR,  
address: \{  
  city : VARCHAR,  
  state : VARCHAR,  
  streetAddress : VARCHAR,  
  state : VARCHAR,  
  postalCode : VARCHAR  
  \},  
phoneNumbers: array(\{  
  aaa : VARCHAR, type : VARCHAR,  
  number : VARCHAR  
  \}),  
cc:array(VARCHAR)
```

See how complex structures are specified using braces ({...}) and arrays using the array keyword and specifying the contents of the array between parentheses.

Also field names can be surrounded by single or double quotes (e.g. "phone number") if they contain white spaces or non-alphanumeric chars.

### Edit Wrapper Parameter values

Enter values for the following wrapper parameters:

Host	<input type="text" value="127.0.0.1"/>
Port	<input type="text" value="27017"/>
	<input type="checkbox"/> Pass-through session credentials
User	<input type="text"/>
Password	<input type="text"/>
Database	<input type="text" value="test"/>
Collection	<input type="text" value="address"/>
Connection String	<input type="text"/>
Fields	<input type="text" value="_id:VARCHAR, firstName:VARCHAR"/>
Introspection query	<input type="text"/>

Mongo DB base View using Fields parameter

### Edit value of 'Fields'

```

_id:VARCHAR,
firstName:VARCHAR
,lastName:VARCHAR,
age:VARCHAR,
address:{city:VARCHAR,state:VARCHAR,streetAddress:VARCHAR,state:VARCHAR,postalCode:VARCHAR},
phoneNumbers:array({aaa:VARCHAR,type:VARCHAR,number:VARCHAR}),
cc:array(VARCHAR)

```

Fields Parameter

## 5.6 CREATE BASE VIEW USING INTROSPECTION QUERY

If you want to not define the schema, you should use the Introspection query parameter. For this you should use a query over the collection, following the syntax for the mongodb method `db.collection.find()`. You should escape the braces.

A example of introspection query is :

```
\{ "properties": \{"LINEARID": "1101121687137", "observed": "", "COUNTYFP": "051", "RTTYP": "M", "FULLNAME": "Airport Rd", "MTFCC": "S1400", "STATEFP": "35"\}\}
```

Edit Wrapper Parameter values

Enter values for the following wrapper parameters:

Host

Port

☐ Pass-through session credentials

User

Password

Database

Collection

Connection String

Fields

Introspection query

roads

mongodb://127.0.0.1:27017/test

\{ "properties": \{"LINEARID": "1101121687137"\}\}

Ok

Cancel

MongoDB Base View edition using Introspection query parameter

## 5.7 CREATE BASE VIEW WITH TLS/SSL

The MongoDB API for Java Driver allows the connection to a MongoDB Server that has TLS/SSL enabled. To create a Base view using SSL, you have to add the option `ssl=true` in the Connection String parameter. `ssl` is one of the [Connection Options of the Connection String URI](#). Alternatively (or additionally), you can also check the Use SSL/TLS checkbox at view creation time.

In addition you have to import the certificate, required to connect to the MongoDB server, into the *TrustStore* of the Java Runtime Environment (JRE) of the Denodo Platform. The section ["Importing the Certificates of Data Sources \(SSL/TLS Connections\)"](#) of the Denodo Platform Installation Guide explains how to do this.



A example of a Base view with TLS/SSL:

Edit Wrapper Parameter values

Enter values for the following wrapper parameters:

Host

Port

☐ Pass-through session credentials

User

Password

Database

Collection

product

Connection String

mongodb://localhost:27017/products?ssl=true

☒ Use SSL/TLS

Fields

Introspection query

Ok

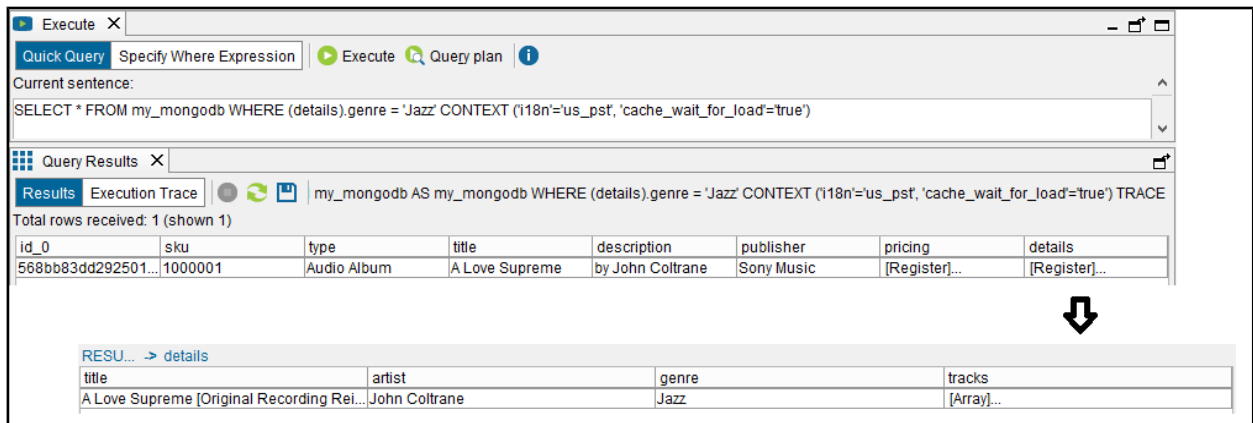
Cancel

MongoDB Base View using SSL

Also note that in order to use **TLSv1.2** in Denodo 6 (with Java 7), you will need to add the property “com.denodo.parser.connection.http.tlsProtocol=TLSv1.2” to the VDBConfiguration.properties file.

## 5.8 QUERYING MONGODB COLLECTIONS

From now on we can query using the VQL Shell tool of the Admin Tool, create derived views, etc.

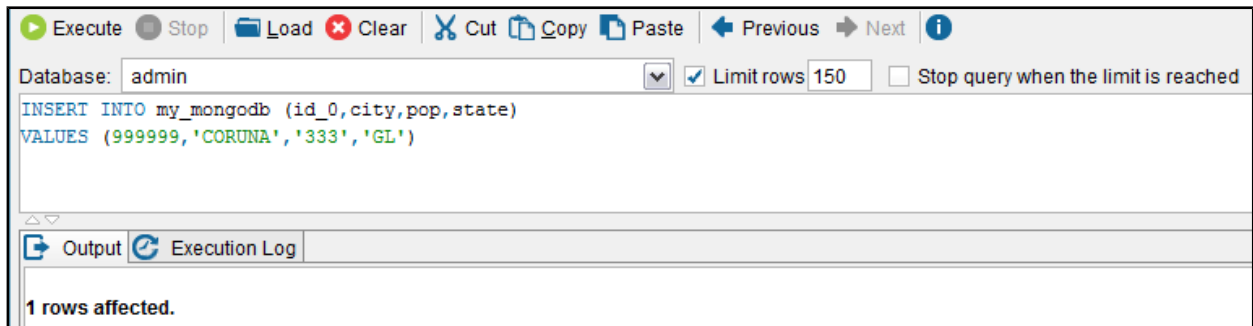


The screenshot shows the Denodo Query Editor interface. The 'Current sentence' field contains a MongoDB query: `SELECT * FROM my_mongodb WHERE (details).genre = 'Jazz' CONTEXT ('i18n'='us_pst', 'cache_wait_for_load'='true')`. The 'Query Results' tab is active, showing a table with 8 columns: `id_0`, `sku`, `type`, `title`, `description`, `publisher`, `pricing`, and `details`. The first row of data is: `568bb83dd292501...`, `1000001`, `Audio Album`, `A Love Supreme`, `by John Coltrane`, `Sony Music`, `[Register]...`, and `[Register]...`. Below the table, a 'details' view is shown, displaying the `title`, `artist`, `genre`, and `tracks` fields for the selected document.

Querying a MongoDB View

## 5.9 INSERT

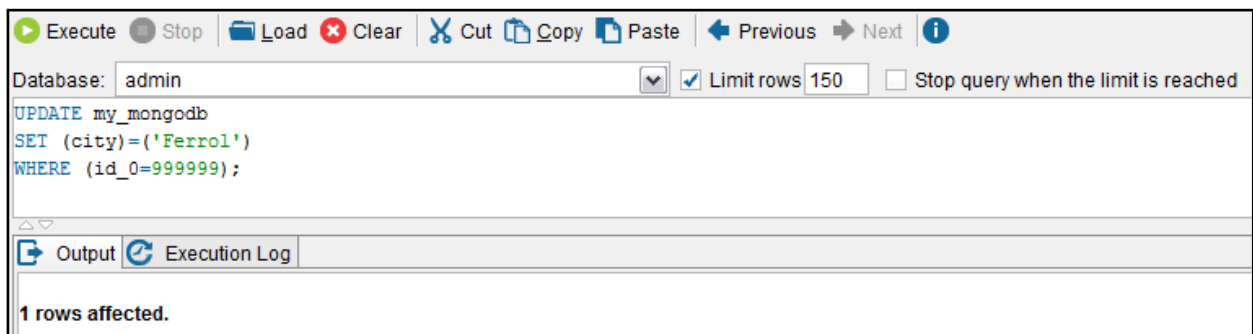
This wrapper allows to make insertion. No complex fields are allowed.



The screenshot shows the Denodo Query Editor interface. The 'Database' dropdown is set to 'admin'. The 'Limit rows' checkbox is checked and set to 150. The 'Stop query when the limit is reached' checkbox is unchecked. The 'Current sentence' field contains an INSERT query: `INSERT INTO my_mongodb (id_0,city,pop,state) VALUES (999999,'CORUNA','333','GL')`. The 'Output' tab is active, showing the message: '1 rows affected.'

## 5.10 UPDATE

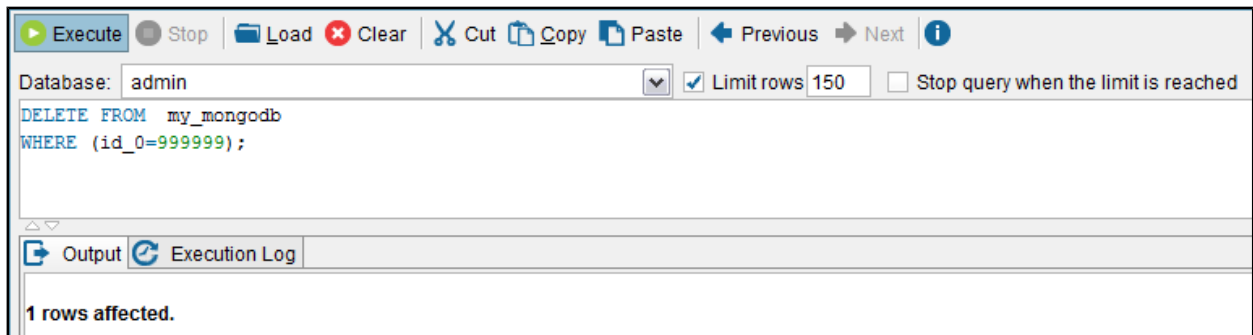
This wrapper allows to make updates in a document. No complex fields are allowed.



The screenshot shows the Denodo Query Editor interface. The 'Database' dropdown is set to 'admin'. The 'Limit rows' checkbox is checked and set to 150. The 'Stop query when the limit is reached' checkbox is unchecked. The 'Current sentence' field contains an UPDATE query: `UPDATE my_mongodb SET (city)='Ferrol' WHERE (id_0=999999);`. The 'Output' tab is active, showing the message: '1 rows affected.'

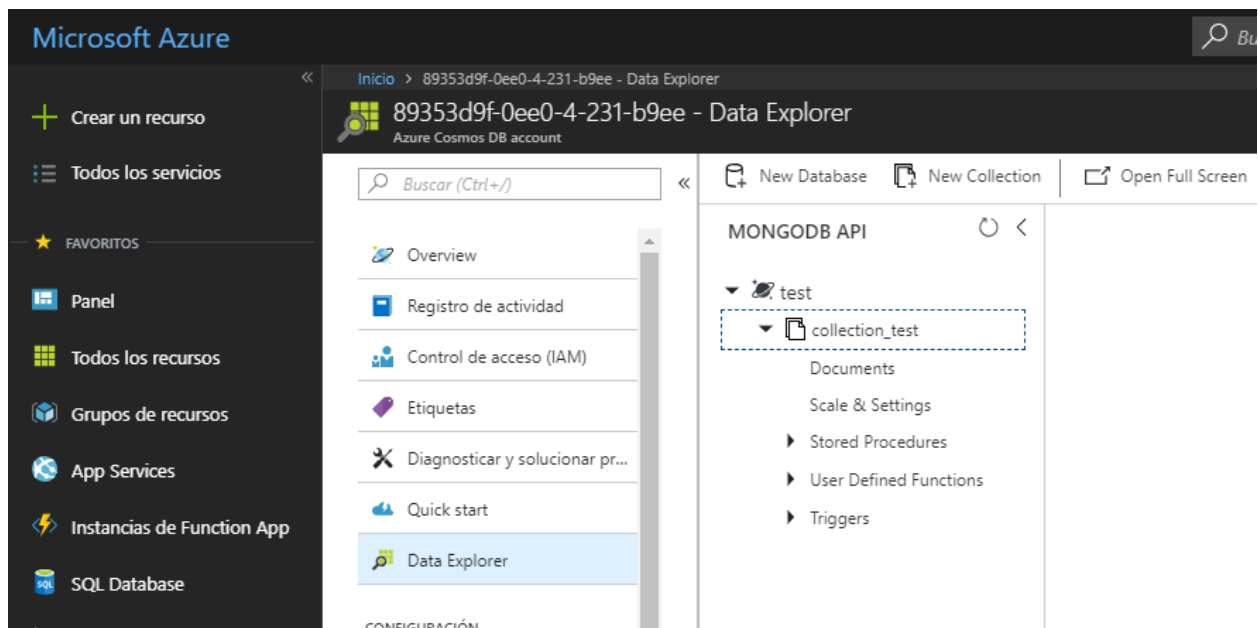
## 5.11 DELETE

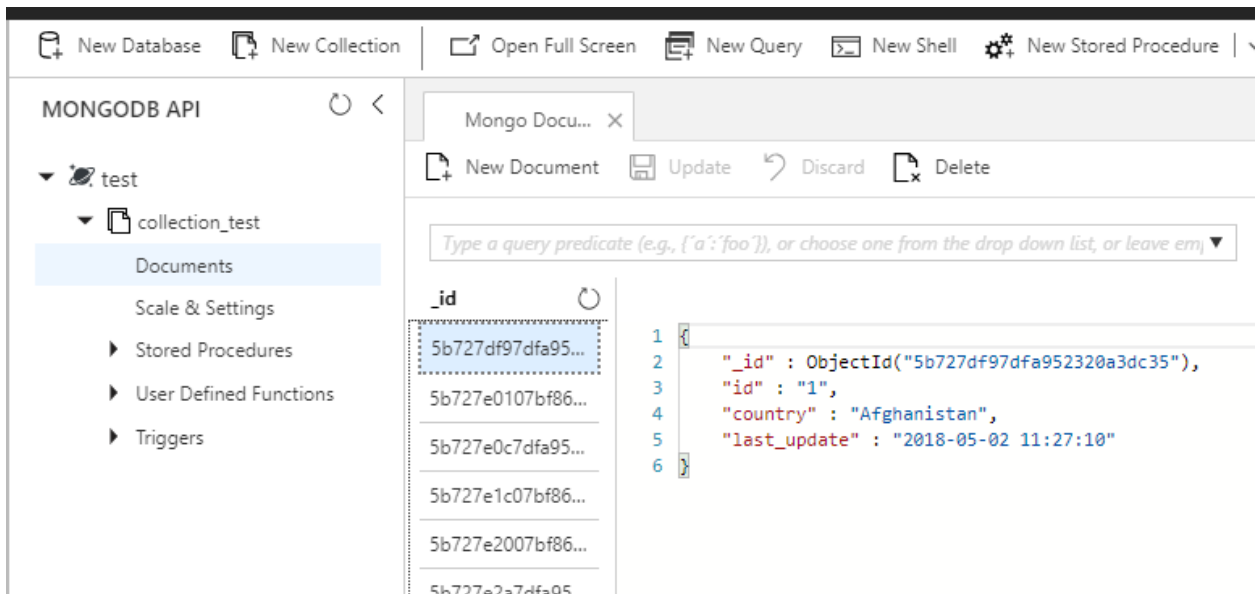
It is allowed to delete documents of a collection.



## 5.12 USING MONGODB WRAPPER TO CONNECT TO AZURE COSMOS DB

You can use this wrapper to access Azure Cosmos DB, taking advantage of their MongoDB protocol API. So, if you have a collection like this one (collection\_test):





You are able to access it creating a new MongoDB base view over a MongoDB data source this way:

1. Create a Custom datasource selecting the denodo-mongodb-customwrapper-7.0 (\*.6.0)
2. Create a base view adding the parameters:
  - a. User: The USERNAME shown in the Connection String section.
  - b. Password: The PRIMARY PASSWORD shown in the Connection String section.
  - c. Collection: The collection we want to introspect.
  - d. Connection String: The connection string with the format `mongodb://{HOST}:{PORT}/{DATABASE}?ssl=true&replicaSet=globaldb`

For our collection\_test example:

Denodo North America & APAC: 525 University Avenue, Suite 31, Palo Alto, CA 94301. USA  
Denodo Iberia & Latino América: Montalbán 5, 28014 Madrid, Spain  
Denodo EMEA: 21<sup>st</sup> Floor, Portland House, Bressenden Place, London SW1E 5RS. UK  
Denodo DACH: Karlstraße 10, 80333 München. Germany

VQL Shell x mongodb.comsodb x mongodb.bv\_comsodb\_country x

Summary Edit Options VQL Execution panel Used by Associations Publish Export

Database: mongodb

View type: Base

View schema:

PK	Field Name	Field Type	Description
	_id	text	
	id	text	
	country	text	
	last_update	text	

Owner: admin Last modifier: admin

Creation: 14-ago-2018 9h 8m 28s Last modification: 14-ago-2018 9h 8m 29s

Swap status: default Cache status: off

Folder: /01 - data sources

Description:

Execute x Query Results x

Results Execution Trace Query: SELECT \* FROM bv\_comsodb\_country CONTEXT ('i18n'='es\_euro', 'cache\_wait\_for\_load'='true')

Total rows received: 10 (shown 10)

_id	id	country	last_update
5b727df97dfa952320a3dc35	1	Afghanistan	2018-05-02 11:27:10
5b727e0107bf8627c04c3b8f	2	Algeria	2006-02-15 03:44:00
5b727e0c7dfa952320a3dc36	3	American Samoa	2006-02-15 03:44:00
5b727e1c07bf8627c04c3b90	4	Angola	2006-02-15 03:44:00
5b727e2007bf8627c04c3b91	5	Anguilla	2006-02-15 03:44:00
5b727e2a7dfa952320a3dc37	6	Argentina	2006-02-15 03:44:00
5b727e5a07bf8627c04c3b92	7	Armenia	2006-02-15 03:44:00
5b727e637dfa952320a3dc38	8	Australia	2006-02-15 03:44:00
5b727e6e7dfa952320a3dc39	9	Austria	2006-02-15 03:44:00
5b727e767dfa952320a3dc3a	10	Azerbaijan	2006-02-15 03:44:00

## 6 LIMITATIONS

---

### 6.1 WRITE OPERATIONS

Create/update operations with complex fields, arrays and registers, on the MongoDB server are not supported through this custom wrapper at its current version.

If you try this operation, you can see the following exception, in the output of the VQL shell:

```
MongoDB wrapper error. Can't find a codec for class  
com.denodo.vdb.engine.customwrapper.value.CustomWrapperStruct.
```

### 6.2 MONGODB BSON TIMESTAMP DATA TYPE

The MongoDB “BSON Timestamp” data type is considered by MongoDB to be for internal use only (see <https://docs.mongodb.com/manual/reference/bson-types/#timestamps>), but is nevertheless supported by this MongoDB custom wrapper by considering values of this type equivalent to MongoDB’s recommended “BSON Date” type.

Note however that “BSON Timestamp” data types are not supported in WHERE clauses in updates/deletes, nor they will be used in INSERT or UPDATE sentences to set new values into MongoDB (new values will be set as BSON Date, not BSON Timestamp). You can use them in select queries, but notice that filtering will be done in memory instead of directly in the data source.

### 6.3 LIMITATIONS IN SELECT SENTENCES

Note that group by operations are not delegated to the MongoDB server, and therefore aggregations are performed in-memory by Virtual DataPort.

Also note that functions are not delegated to the mongoDB server, and therefore any WHERE clauses containing functions will be executed by Virtual DataPort after retrieving the target data, and not at the mongoDB server.